

Emo-DiVA: Emotion Disentanglement via Variational Autoencoders

Tebe
Nigrelli

Federico
Pezzoli

Alessandro
Reali

Luca
Ricci

Ali Emre
Senel

Abstract

Encoder-decoder transformer models do not have interpretable hidden representations. Addressing this limitation can help with model alignment: disentangling the inner representations of different information can make it easier to access, edit and use. Current approaches use Variational Autoencoders (VAEs) to produce latents with interpretable components, with limited success. We experiment with a T5 + FactorVAE architecture, progressively expanding it to separate emotion from context. We test the model on classification and translation of emotion. Classification performs better than the BERT baseline, while translation fails despite our additions. The results suggest that we achieve partial disentanglement of emotion, but architecture and dataset are jointly insufficient to implement translation.

<https://github.com/tebe-nigrelli/NLP-Latent-Learning>

1 Introduction

Pretrained text-to-text encoder-decoder Transformer models such as T5 and FLAN-T5 learn powerful linguistic representations [1], [2], but their internal states are not exposed as latents that users can inspect, edit, or constrain. This limits explainability and alignment, since different aspects of meaning often remain entangled in the hidden representation. Recent work addresses this issue by processing the hidden states of a pretrained sequence-to-sequence model using a Variational Autoencoder [3]. This mechanism forces the model to map text into structured latent variables before decoding them back into language.

We apply this framework by injecting a FactorVAE [4] between the encoder and decoder of FLAN-T5, with the goal of separating emotional information from contextual content in the latent space. We evaluate the resulting model on emotion classification, latent-space disentanglement, sentence reconstruction, and emotion translation. By reconstruction, we mean the model’s ability to reproduce the input sentence when the emotion-related latent component is left unchanged. By translation, we mean modifying the emotional tone of a sentence while preserving its contextual meaning, as in Figure 1. We outperform the BERT baseline in emotion




Figure 1: Emotion translation example.

classification, investigate emotion disentanglement, and achieve reconstruction, but not translation.

2 Experimental Setup

Backbone. We use FLAN-T5-base as our base model because it was pretrained on a large dataset, and its weights are publicly available. Also, its encoder-decoder structure allows us to train the entire architecture. Its limited size also enables faster training, thus more experiments [2].

VAE. We use a FactorVAE because it encourages representations with uncorrelated factors and supports generation from the learned latent space [4]. We split latent vectors into two subspaces: emotion and context.

Dataset. We use GoEmotions, a corpus of 54k English Reddit samples annotated with 27 emotion labels plus 1 for neutral. We select it for its broad coverage, size, multi-label format, and canonical train-validation-test split, which lets us compare our classifier with published results [5].

2.1 Tasks and Evaluation

Disentanglement is difficult to evaluate intuitively, as it involves high-dimensional spaces and a complex structure for meaning. We evaluate separation using metrics that capture different properties of the latent space, and from two downstream tasks: **emotion classification** and **emotion transla-**

tion. Since no models achieved meaningful results in emotion translation, we also evaluate reconstruction as a diagnostic intermediate step.

The main disentanglement metrics are: **mutual information gap** (MIG), where higher values indicate stronger separation between emotion and context information; **emotion R^2** , which measures how much emotion-label variance can be predicted from the emotion latent using linear regression; and **separation R^2** , defined as the difference between emotion R^2 and the largest non-emotion R^2 , which may come from the context latent or residual connections depending on the architecture. MIG and emotion R^2 range from 0 to 1, with higher values being better. Separation R^2 is also better when larger, but unless clipped or rescaled, it can range from -1 to 1.

We evaluate classification with **macro-F1** and **micro-F1**. Reconstruction is measured with **exact match** and **token F1**. For translation, we only report fluency through **GPT-2 perplexity** (PPL), since outputs are often nonsensical. We also report **median-PPL**, because **mean-PPL** suffers from the presence of large outliers.

3 Exploratory Data Analysis

We analyze label frequency, number of emotions per input, input length, and label co-occurrence on 54263 GoEmotions comments. Excluding neutral, the three most frequent classes are admiration (5122), approval (3687), and gratitude (3372), while grief (96), pride (142), relief (182), and nervousness (208) are rare. The dataset is therefore strongly imbalanced. Most inputs contain a single emotion: 83.8% of examples have one label, 15.0% have two, 1.2% have three, 0.07% have four, and fewer than 0.01% have five. This confirms that the task is mostly single-label, but still requires sigmoid outputs instead of a softmax classifier. Inputs are short: the median length is 12 words, although the character distribution has a long tail up to 703 characters. These findings motivate reporting macro-




Figure 2: Label frequency in the dataset




Figure 3: 1 → base; 2 → residual bottleneck; 3 → LoRA on T5 decoder + copy loss; 4 → adversarial classifier.

F_1 together with micro- F_1 , since micro- F_1 can hide poor performance on rare emotions.

4 Architectures and Results

Architecture 1. Our base architecture (blue in Figure 3) inserts a FactorVAE between the T5 encoder and decoder. The FactorVAE encoder maps transformer hidden states into a sequence of latent representations, *separating emotion and context branches*; the FactorVAE decoder then projects these latent representations back to the hidden-state dimension as T5 decoder memory. Depending on the task, the T5 decoder finally reconstructs or translates the input, producing a new sequence of tokens.

For classification, we pool the sequence of emotion latents using attention weights generated from the T5 encoder hidden states. Pooling can be done either **jointly**, across the emotion subspace, or independently **per emotion factor**. The resulting representation is then passed to an MLP classifier. Similar to pooling, the classifier can either use a **single classification head** or use **separate heads** for individual latent dimensions. For a size comparison, see Table 2.

We train our model using a combination of different losses: a standard binary classification loss, a reconstruction loss, a KL loss for the VAE, and a total-correlation loss, specific for the disentanglement in the FactorVAE [4]. The respective weights are $\alpha_{cl} = 4.0$, $\alpha_{rec} = 2.0$, $\alpha_{KL} = 0.05$, $\alpha_{TC} = 0.1$.

We test this architecture by ablating on the dimensionality of the emotion and context latents, the pooling mode, and the classifier mode. Since our dataset contains 28 labels, we set the emotion latent dimensionality to either 28 or 56, while for the context latent we test 64 and 128 dimensions Table 1.

The results show that joint pooling favors disentanglement, especially in terms of separation R^2 ($-0.003 \rightarrow 0.009$). A lower-dimensional context latent has a similar effect on separation R^2 ($-0.017 \rightarrow 0.023$). Moreover, a 56-dimensional emotion latent is better separated from the context latent than a 28-dimensional emotion latent ($-0.007 \rightarrow 0.013$). Branch MIG follows the same trend: averaged across ablations, joint pooling improves over per-emotion pooling ($0.227 \rightarrow 0.195$), a 64-dimensional context latent improves it over a 128-dimensional context latent ($0.223 \rightarrow 0.199$), and a 56-dimensional emotion latent improves it over a 28-dimensional one ($0.220 \rightarrow 0.203$). The best Branch MIG value (0.246) is obtained by the same configuration that maximizes separation R^2 Table 1.

Classification achieves on average **~ 0.5 macro-F1**: since it remains above the BERT baseline (**0.46 macro-F1**, as per [5]) consistently across ablations, we keep the 56-dimensional emotion latent. Consequently, we configure the MLP with **one classifier head for every two emotion latent dimensions**, since this classifier mode achieves the best results in both micro- and macro-F1. Reconstruction and translation fail for all tested setups.

Architecture 2. We add a **residual connection** (red in Figure 3) to route features around the VAE. The residual connection is a small MLP bottleneck over T5 hidden states: a linear layer reducing dimension from 768 to 64; GELU activation; 50% dropout; a final linear layer mapping the hidden states back to their original dimension. The goal is to stabilize training and avoid catastrophic forgetting of the T5 decoder. Although classification performance matches Architecture 1, **reconstruction remains poor**, with a token-F1 value of 0.1. However, skip connections do not hurt disentanglement, and in fact they improve it: residual R^2 is considerably low (< 0.01) compared to the context R^2 (< 0.1 , an improvement compared to Architecture 1) and, more importantly, emotion R^2 (≈ 0.15).

Architecture 3. Since the VAE’s reconstruction loss is too weak to teach the decoder how to read the new decoder memory, we use **LoRA** adapters [6] to train the T5 decoder efficiently and add a token-level **copy loss** (brown section in Figure 3). Copy loss is a teacher-forced cross-entropy loss of the T5 decoder, forcing the output from decoder memory M to match the input:

$$L_{\text{copy}} = -\frac{1}{T} \sum_{t=1}^T \log p_{\theta}(x_t \mid x_{<t}, M).$$

This model preserves the previous disentanglement results, while slightly improving classification, increasing micro-F1 to 0.597 and macro-F1 to 0.517, and having LoRA rank = 16 and LoRA $\alpha = 32$. As intended, copy loss boosts reconstruction to consistently good results, with 90% perfect matches across samples and **$\sim 99\%$ token-F1**. Translation becomes more fluent, but still fails, with a median-PPL of 96 and a copy rate of almost 80%. Since modifying the emotion vector does not influence reconstruction, we believe the model **stores emotional information in the context vector**.

Architecture 4. We add two **adversarial MLPs**, one on the context latents and one on the residual branch. Each MLP pools its input and learns to predict emotion. At the same time, the MLPs send inverse gradients backward, which push emotion out of context, **increasing separation R^2 to 0.1**. Surprisingly, the adversarial components are not enough to shake translation from mere reconstruction, leading to similar results, which suggests the **model overfits on reconstruction**, using decoder-ready context and residuals.

Architecture 5. Since residual and adversarial variants show that regularization alone does not solve the shortcut problem, we rethink the architecture. We replace the residual connection and VAE decoder with a split decoder and introduce pooled emotion injection, inspired by feature-wise linear modulation (FiLM) [7] and emotion modulation in LLM-TTS [8]. The split decoder maps context latents to token memory M_c and emotion latents to token memory M_e . The pooled emotion summary produces a general bias b_e and FiLM parameters $\Delta\gamma_e$ and β_e . The decoder memory is computed as:

$$M = (1 + \tanh(\Delta\gamma_e)) \cdot (M_c + M_e + b_e) + \beta_e.$$

During training, inspired by Park et al. [3], we add a warmup schedule to the FactorVAE. The KL and TC loss weights start at zero, allowing the split-FiLM decoder to learn reconstruction before regularization. The weights are then gradually increased over the training. We evaluate how different schedules impact the model Table 3, varying the number of epochs (0, 2, 4, 6) where KL and TC losses are deactivated, followed by four warm-up epochs. The best configuration, with 4 zero-epochs, improves classification, reaching 59% micro-F1 and **52% macro-F1**, and improves disentanglement, in particular increasing **separation R^2 above 0.14**. As for translation, the model shows modest improvements, reducing copy rate down to 65%, without




Figure 4: 5 → split-FiLM VAE Decoder.

affecting median-PPL; but still fails to achieve emotion translation.

Reproducibility.

We tune the model by varying the latent bottleneck and the components that most affect task performance, while keeping the optimization setup fixed across comparable runs. We use learning rate 2×10^{-4} , train batch size 8, evaluation batch size 16, weight decay 10^{-2} , 10 training epochs, and 0.1 warmup ratio. This setup is stable enough for our experiments, so we focus the search on scalar emotion factors, context-vector size, pooling strategy, skip connections, decoder LoRA, adversarial regularization, and KL/TC warmup. The best trade-off uses a compact 64-dimensional context vector, with micro- F_1 around 0.59 and macro- F_1 around 0.52. The repository includes a *reproducible_scripts/* folder with the Slurm scripts used to train, evaluate, and summarize the reported runs: *train_eval**, *summarize**, *run_reconstruction**, *run_translation**, and *collect_beta_classification_samples.slurm*. Attention analyses are also reproduced using *analyze_attention_maps.py*, *visualize_attention_examples.py*, and *aggregate_attention_words_by_emotion.py*. Further architectural details and ablations are reported in Table 2.

5 Discussion

The final architecture surpasses BERT-level accuracy in emotion classification [5], [9] and successfully reconstructs sentences, but fails at emotion translation. More specifically, translation experiments show that the decoder rewrites the input with small lexical edits at times, generating fluent but semantically weak rewrites in other cases. In fact, the model never achieves proper emotion translations. However, the model captures emotion well enough for recognition, but not well enough to manipulate it independently from context.

For example, using as input the sentence “I’m really sorry about your situation :(Although I love the names Sapphira, Cirilla, and Scarlett!”, labeled in the dataset as “remorse” and “sadness”, our model

correctly labels it and reconstructs it by changing the second name in the list. However, translation to “joy” fails, and outputs the same sentence.

This inability indicates that the model **does not yet have a stable mechanism** for moving from one emotional region of the latent space to another while preserving sentence context. However, the architecture **can organize emotion information** into a classifiable subspace. This aspect is not surprising, since the FactorVAE forces the latent space of the T5 model into a smaller space, where the classification MLP can successfully recover the information for the emotion recognition task. The success of sentence reconstruction can also be explained by the introduction of the copy loss from architecture 3, which encourages the model to correctly reproduce the input.

The failure in the emotion translation task, instead, can be attributed to multiple reasons. GoEmotions is a small dataset when compared to the training data of the T5 model. Moreover, it provides emotion labels, without parallel sentence pairs, showing the same context with different emotions. Also, **emotion and context are often deeply entangled in natural language**, so changing the emotion without changing the semantic content is intrinsically difficult.

Copy loss may have influenced the outcome of the task by pushing the model toward the total reconstruction too strongly, introducing a **bias toward copying the input** instead of modifying emotional content. Nonetheless, the small lexical edits and the few semantically weak rewrites obtained by our model suggest that we are moving in the correct direction.

We also visualized attention weights, confirming the model attends to plausible lexical cues Table 4, and the per-label most attended words are sensible Table 5. We see that the latent pooling attention is most interpretable when the emotion has a conventional lexical realization. Emotions attend to typically relevant words, such as “gratitude” attending to *thanks* Table 5. Also, KL divergence over the phrases shows the model spreads attention over the tokens, as opposed to focusing only on easy-to-spot single words, which would appear as a concentrated attention map Section B.3.

6 Related Work

Fu et al. [10] perform *style transfer* with non-parallel corpora by disentangling content and style through adversarial training. They evaluate both transfer strength and content preservation; we consider a

stricter setting, modifying emotion while preserving contextual meaning.

Park and Lee [3] convert pretrained encoder-decoder transformers into text VAEs using two-stage finetuning: first autoencoding with zero KL weight, followed by gradual KL warmup. We adopt only the warmup strategy as stabilizer in Architecture 5.

Perez et al. [7] introduce FiLM, which modulates intermediate features through feature-wise affine transformations based on conditioning information. While they condition visual reasoning on language, we use FiLM-style modulation in Architecture 5 to inject pooled emotion information into T5 decoder memory.

7 Conclusion

We investigated whether the latent space of an encoder-decoder model can be disentangled into emotion and semantic context. To this end, we inserted a FactorVAE between the encoder and decoder of a T5 model and refined the architecture toward a FiLM-style design. The final models achieved competitive performance on emotion classification and sentence reconstruction, but failed to perform reliable emotion translation. These results suggest that the proposed research direction may be limited by dataset size, model choice, and an unavoidable emotional entanglement in the latent representation. Future work could explore alternative encoder-decoder models, such as BART, which is pretrained as a denoising autoencoder. Moreover, weakening copy loss through a stochastic mask may help reduce overfitting. Finally, applying dimensionality reduction and latent-space visualization techniques can help clarify why disentanglement and emotion-controlled generation remain challenging.

8 Limitations

Dataset. We only used the GoEmotions dataset, which is small compared to the 750 GB C4 corpus used for training FLAN-T5-base [1]. Size limits lexical and emotional variation, by not providing a wide coverage of both emotion and context independently. Multiple sources of bias also exist in the Reddit dataset: it is almost completely in English, it contains community-specific patterns, and was labelled by human annotators.

Model. All our results were only tested with FLAN-T5-base as backbone. Although we expect our results to carry to other architectures, we only tested the FLAN-T5-base pretrained model, which limits our findings.

Ablations. All ablation settings were run once. Consequently, performance metrics may be affected by poor initialization or randomness, although we do not see evidence in our experiments to support this claim.

Bibliography

- [1] C. Raffel *et al.*, “Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer,” 2020, pp. 1–67. [Online]. Available: <http://jmlr.org/papers/v21/20-074.html>
- [2] H. W. Chung *et al.*, “Scaling Instruction-Finetuned Language Models,” 2022. [Online]. Available: <https://arxiv.org/abs/2210.11416>
- [3] S. Park and J. Lee, “Finetuning Pretrained Transformers into Variational Autoencoders,” 2021. [Online]. Available: <https://arxiv.org/abs/2108.02446>
- [4] H. Kim and A. Mnih, “Disentangling by Factorising,” 2019. [Online]. Available: <https://arxiv.org/abs/1802.05983>
- [5] D. Demszky, D. Movshovitz-Attias, J. Ko, A. Cowen, G. Nemade, and S. Ravi, “GoEmotions: A Dataset of Fine-Grained Emotions,” in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, D. Jurafsky, J. Chai, N. Schluter, and J. Tetreault, Eds., Online: Association for Computational Linguistics, Jul. 2020, pp. 4040–4054. doi: [10.18653/v1/2020.acl-main.372](https://doi.org/10.18653/v1/2020.acl-main.372).
- [6] E. J. Hu *et al.*, “LoRA: Low-Rank Adaptation of Large Language Models,” 2021. [Online]. Available: <https://arxiv.org/abs/2106.09685>
- [7] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, “FiLM: Visual Reasoning with a General Conditioning Layer,” 2017. [Online]. Available: <https://arxiv.org/abs/1709.07871>
- [8] S. Wang, A. Chen, and T. Zhao, “Beyond Global Emotion: Fine-Grained Emotional Speech Synthesis with Dynamic Word-Level Modulation,” 2025. [Online]. Available: <https://arxiv.org/abs/2509.20378>
- [9] A. Harutyunyan and S. Kumar, “Fine-Grained Emotion Detection on GoEmotions: Experimental Comparison of Classical Machine Learning, BiLSTM, and Transformer Models,” 2026. [Online]. Available: <https://arxiv.org/abs/2601.18162>
- [10] Z. Fu, X. Tan, N. Peng, D. Zhao, and R. Yan, “Style Transfer in Text: Exploration and Evaluation,” 2017. [Online]. Available: <https://arxiv.org/abs/1711.06861>

A Architecture 1 Ablations

The table compares Architecture 1 across pooling strategies, emotion-factor counts, and context dimensions. The strongest overall configuration uses 56 emotion factors. Different settings achieve the best trade-offs across classification, emotion prediction, separation, and branch disentanglement.

Pooling mode	Emo	Context	micro-F1	macro-F1	Emo R ²	Sep R ²	B. MIG
joint_emotion_context	28	64	0.580	0.508	0.149	0.019	0.241
joint_emotion_context	28	128	0.579	0.514	0.151	-0.017	0.202
per_emotion_dim	28	64	0.576	0.508	0.131	0.002	0.195
per_emotion_dim	28	128	0.580	0.491	0.130	-0.032	0.172
joint_emotion_context	56	64	0.585	0.504	0.164	0.041	0.246
joint_emotion_context	56	128	0.587	0.514	0.166	-0.006	0.218
per_emotion_dim	56	64	0.591	0.504	0.155	0.030	0.211
per_emotion_dim	56	128	0.589	0.505	0.155	-0.012	0.203

Table 1: Architecture 1 results

We report architecture size for the different ablations of Architecture 1, depending on latent dimensions, pooling mode and the number of classifier heads. In the table, B is batch size and T is input sequence length. Note that T5 hidden states have dimension 768.

Emo	Ctx	Pooling mode	Pooling shape [B,T,768] → ...	Class. heads	Class. shape
28	64	per-factor	separate emotion pools → [B,28,8]	multi	28 × MLP(8 → 32 → 1)
28	64	joint	shared joint evidence → [B,28,8]	multi	28 × MLP(8 → 32 → 1)
28	128	per-factor	28 separate emotion pools → [B,28,8]	multi	28 × MLP(8 → 32 → 1)
28	128	joint	shared joint evidence → [B,28,8]	multi	28 × MLP(8 → 32 → 1)
56	64	per-factor	56 separate emotion pools → [B,56,8]	multi	28 × MLP(16 → 32 → 1)
56	64	joint	shared joint evidence → [B,56,8]	multi	28 × MLP(16 → 32 → 1)
56	128	per-factor	56 separate emotion pools → [B,56,8]	multi	28 × MLP(16 → 32 → 1)
56	128	joint	shared joint evidence → [B,56,8]	multi	28 × MLP(16 → 32 → 1)

Table 2: Architecture 1 size comparison

B Architecture 5

B.1 Ablation Results

Epoch	micro-F1	macro-F1	Sep R ²	Emo R ²	B. MIG
0	0.5880	0.5200	0.1147	0.1411	0.2430
2	0.5924	0.5193	0.1354	0.1456	0.2427
4	0.5914	0.5240	0.1414	0.1486	0.2516
6	0.5822	0.5196	0.1422	0.1543	0.2563

Table 3: KL warmup start epoch. Higher warmup start improves all three disentanglement metrics at a small classification cost.

B.2 Attention Maps

Following multi-hot classification, the model assigns each label an independent probability $\hat{p}$. A calibration set is used to identify the threshold value. A label is predicted as active if its score exceeds the label-specific threshold θ , i.e. $\hat{p} \geq \theta$.

Labels	Predicted	$\hat{p}$	θ	Phrase token importance - Average Attention Map	Comment
gratitude	gratitude	0.999	0.921	I didn't know that thank you for teaching me 0.03 0.09 0.01 0.09 0.13 0.02 0.01 0.34 0.03 0.06 0.09 0.02 something today ! 0.04 0.02 0.04	The highest mass falls on the explicit cue "thank".
amusement	amusement	0.997	0.921	Lo I ! But I love your last name though X D 0.18 0.23 0.03 0.05 0.02 0.28 0.03 0.01 0.01 0.03 0.02 0.01 0.02 0.08	The map peaks on love, but assigns mass to "Lol" fragments.
love	love	0.965	0.921		
—	approval	0.884	0.804		
realization	—	0.403	0.637	I had to write in the Swan gas . See m s far better 0.02 0.06 0.05 0.04 0.03 0.02 0.02 0.02 0.04 0.08 0.11 0.02 0.08 0.26 than the choices given 0.05 0.02 0.04 0.02 0.03	Attention concentrates on "better", supporting approval, while missing the gold realization label.

Table 4: Score cell is green when the score exceeds the threshold, showing classification, and red otherwise. Label and predicted-emotion cells are highlighted when prediction and gold labels disagree.

Emotion	Word	Mass
gratitude	thanks	0.213
	thank	0.182
	you	0.017
	appreciate	0.013
	welcome	0.009
love	love	0.352
	like	0.046
	loved	0.019
	favorite	0.009
	loving	0.007
realization	realize	0.031
	realized	0.029
	wow	0.023
	forgot	0.022
	never	0.019

Table 5: Top attended words by emotion.

B.3 Attention Weights Entropy Analysis

The quantile table suggests that attention concentration differs systematically by predicted emotion. Lexically explicit emotions such as *gratitude*, *love*, and *remorse* show relatively low entropy, indicating sharp phrase-level evidence around cue words. *Gratitude* also has the largest mean head-template KL, suggesting stronger head-level variation around these cues. In contrast, *neutral* has the lowest KL and highest entropy, indicating broad, shared attention rather than a localized emotional cue. More discourse-dependent labels such as *approval*, *realization*, *confusion*, and *curiosity* are also relatively high-entropy, consistent with the need to interpret context or stance rather than isolated keywords.

Emotion	n	KL Divergence Value				H Information Entropy Value			
		μ	Q1	50	Q3	μ	Q1	50	Q3
gratitude	329	0.167	0.128	0.168	0.207	0.748	0.715	0.764	0.794
amusement	317	0.146	0.106	0.139	0.176	0.809	0.768	0.817	0.852
desire	37	0.144	0.108	0.141	0.175	0.760	0.714	0.777	0.820
curiosity	459	0.143	0.115	0.140	0.169	0.870	0.840	0.887	0.910
optimism	184	0.143	0.102	0.141	0.173	0.825	0.781	0.823	0.875
embarrassment	24	0.140	0.112	0.132	0.159	0.810	0.749	0.809	0.844
confusion	181	0.140	0.110	0.137	0.164	0.864	0.834	0.877	0.906
remorse	97	0.137	0.108	0.135	0.164	0.747	0.692	0.760	0.806
fear	106	0.129	0.098	0.124	0.156	0.806	0.771	0.819	0.852
sadness	276	0.128	0.095	0.122	0.157	0.807	0.750	0.824	0.868
disgust	187	0.127	0.097	0.116	0.155	0.831	0.793	0.831	0.873
annoyance	286	0.121	0.097	0.117	0.147	0.849	0.813	0.860	0.898
love	262	0.120	0.091	0.114	0.148	0.742	0.694	0.751	0.814
anger	189	0.119	0.095	0.115	0.146	0.845	0.809	0.861	0.906
realization	179	0.117	0.078	0.112	0.142	0.858	0.831	0.868	0.904
disappointment	259	0.115	0.088	0.110	0.140	0.831	0.790	0.848	0.881
grief	28	0.113	0.085	0.103	0.149	0.848	0.825	0.854	0.876
admiration	504	0.113	0.076	0.108	0.141	0.812	0.769	0.816	0.860
joy	221	0.111	0.078	0.108	0.142	0.768	0.726	0.773	0.825
surprise	165	0.110	0.069	0.103	0.140	0.809	0.768	0.822	0.855
caring	287	0.107	0.075	0.101	0.134	0.846	0.797	0.859	0.907
disapproval	351	0.106	0.083	0.102	0.128	0.859	0.833	0.871	0.897
pride	10	0.104	0.086	0.107	0.111	0.773	0.743	0.758	0.796
relief	10	0.099	0.081	0.087	0.105	0.740	0.677	0.774	0.834
nervousness	19	0.099	0.084	0.100	0.108	0.784	0.726	0.778	0.834
excitement	170	0.095	0.058	0.091	0.126	0.813	0.761	0.819	0.859
approval	525	0.094	0.066	0.090	0.120	0.862	0.830	0.870	0.905
neutral	2627	0.088	0.061	0.084	0.109	0.899	0.875	0.906	0.933

Table 6: KL and normalized-entropy distributions of the attention pooling weights by predicted-positive emotion. Quartile columns report Q1 / median / Q3

B.4 Translation Samples

Architecture 5 emotion translation samples: in some cases the model leaves the text unchanged, even when source emotion conflicts with the target. Otherwise, the model edits the sentence in ways unrelated to (or opposite of) the target emotion: eg. substituting *hardly* with *sadly* for a joy target, and deleting a middle clause without emotional restructuring.

Target	Input	Output
neutral	Girlfriend weak as well, that jump was pathetic.	Girlfriend weak as well, that jump was pathetic.
sadness	Lol! But I love your last name though. XD	Lol! But I love your last name though. XD
joy	I never wanted to punch osap harder after seeing that However not too hardly I cant afford them taking everything away	I never wanted to punch osap harder after seeing that However not too sadly I cant afford them taking everything away
sadness	What's that like? Like what's the thought process? I dunno. I know what's a weird question..i just can't imagine	What's that like? Like what's a weird question?..i just can't imagine

Table 7: Translation (Failure) Examples